

OBJECTIVE :- Apply different MLP Architectures on MNIST dataset

In [1]:

```
# Importing Libraries
from keras.utils import np_utils
from keras.datasets import mnist
import seaborn as sns
from keras.initializers import RandomNormal
import matplotlib.pyplot as plt
%matplotlib inline
import numpy as np
import time

# the data, shuffled and split between train and test sets
(X_train, Y_train), (X_test, Y_test) = mnist.load_data()

print("Number of training examples :", X_train.shape[0], "and each image is of shape (%d, %d)" % (X_train.shape[1], X_train.shape[2]))
print("Number of test examples :", X_test.shape[0], "and each image is of shape (%d, %d)" % (X_test.shape[1], X_test.shape[2]))
```

Using TensorFlow backend.

Downloading data from <https://s3.amazonaws.com/img-datasets/mnist.npz> (<http://s3.amazonaws.com/img-datasets/mnist.npz>)

11493376/11490434 [=====] - 1s 0us/step

Number of training examples : 60000 and each image is of shape (28, 28)

Number of test examples : 10000 and each image is of shape (28, 28)

In [2]:

```
# if you observe the input shape its 3 dimensional vector
# for each image we have a (28*28) vector
# we will convert the (28*28) vector into single dimensional vector of 1 * 784

X_train = X_train.reshape(X_train.shape[0], X_train.shape[1]*X_train.shape[2])
X_test = X_test.reshape(X_test.shape[0], X_test.shape[1]*X_test.shape[2])

# after converting the input images from 3d to 2d vectors

print("Number of training examples :", X_train.shape[0], "and each image is of shape (%d)" % (X_train.shape[1]))
print("Number of test examples :", X_test.shape[0], "and each image is of shape (%d)" % (X_test.shape[1]))
```

Number of training examples : 60000 and each image is of shape (784)

Number of test examples : 10000 and each image is of shape (784)

In [3]:

```
# An example data point
print(X_train[0])
```

```
[ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  3  18  18  18 126 136 175  26 166 255
247 127  0  0  0  0  0  0  0  0  0  0  0  0  30  36  94 154
170 253 253 253 253 253 225 172 253 242 195  64  0  0  0  0  0  0
  0  0  0  0  0  49 238 253 253 253 253 253 253 253 251  93  82
 82  56  39  0  0  0  0  0  0  0  0  0  0  0  0  18 219 253
253 253 253 253 198 182 247 241  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  80 156 107 253 253 205  11  0  43 154
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  14  1 154 253  90  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  139 253 190  2  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  11 190 253  70  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  35 241
225 160 108  1  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  81 240 253 253 119  25  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  45 186 253 253 150  27  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  16  93 252 253 187
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  249 253 249  64  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  46 130 183 253
253 207  2  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  39 148 229 253 253 253 250 182  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  24 114 221 253 253 253
253 201  78  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  23  66 213 253 253 253 253 198  81  2  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  18 171 219 253 253 253 253 195
 80  9  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
 55 172 226 253 253 253 253 244 133  11  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0 136 253 253 253 212 135 132  16
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0]
```


In [0]:

```
# this function is used draw Categorical Crossentropy Loss VS No. of epochs plot
def plt_dynamic(x, vy, ty):
    plt.figure(figsize=(10,5))
    plt.plot(x, vy, 'b', label="Validation Loss")
    plt.plot(x, ty, 'r', label="Train Loss")
    plt.xlabel('Epochs')
    plt.ylabel('Categorical Crossentropy Loss')
    plt.title('\nCategorical Crossentropy Loss VS Epochs')
    plt.legend()
    plt.grid()
    plt.show()
```

(1). Softmax Classifier with 2 hidden layers

(1.a) Without dropout and Batch normalization

In [8]:

```

from keras.models import Sequential
from keras.layers import Dense, Activation
from keras.initializers import he_normal

# some model parameters

output_dim = 10
input_dim = X_train.shape[1]

batch_size = 128
nb_epoch = 20

# Initialising model
model_2 = Sequential()

# Adding first hidden layer
model_2.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=he_r

# Adding second hidden layer
model_2.add(Dense(256, activation='relu', kernel_initializer=he_normal(seed=None)))

# Adding output layer
model_2.add(Dense(output_dim, activation='softmax'))

# Printing model Summary
print("Model Summary :- \n",model_2.summary())

# Compiling the model
model_2.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

# Fitting the data to the model
history_2 = model_2.fit(X_train, y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1

```

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow/python/framework/op_def_library.py:263: colocate_with (from tensorflow.python.framework.ops) is deprecated and will be removed in a future version.
Instructions for updating:
Colocations handled automatically by placer.

Layer (type)	Output Shape	Param #
dense_1 (Dense)	(None, 512)	401920
dense_2 (Dense)	(None, 256)	131328
dense_3 (Dense)	(None, 10)	2570
Total params: 535,818		
Trainable params: 535,818		
Non-trainable params: 0		

Model Summary :-

None
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow/python/ops/math_ops.py:3066: to_int32 (from tensorflow.python.ops.math_ops) is deprecated and will be removed in a future version.
Instructions for updating:
Use tf.cast instead.
Train on 60000 samples, validate on 10000 samples

```
Epoch 1/20
60000/60000 [=====] - 8s 132us/step - loss: 0.8276
- acc: 0.7720 - val_loss: 0.3774 - val_acc: 0.8896
Epoch 2/20
60000/60000 [=====] - 7s 123us/step - loss: 0.3455
- acc: 0.8997 - val_loss: 0.3008 - val_acc: 0.9122
Epoch 3/20
60000/60000 [=====] - 7s 121us/step - loss: 0.2907
- acc: 0.9147 - val_loss: 0.2613 - val_acc: 0.9213
Epoch 4/20
60000/60000 [=====] - 7s 120us/step - loss: 0.2510
- acc: 0.9266 - val_loss: 0.2327 - val_acc: 0.9303
Epoch 5/20
60000/60000 [=====] - 7s 120us/step - loss: 0.2174
- acc: 0.9358 - val_loss: 0.2067 - val_acc: 0.9391
Epoch 6/20
60000/60000 [=====] - 7s 120us/step - loss: 0.1882
- acc: 0.9445 - val_loss: 0.1868 - val_acc: 0.9446
Epoch 7/20
60000/60000 [=====] - 7s 120us/step - loss: 0.1647
- acc: 0.9520 - val_loss: 0.1594 - val_acc: 0.9523
Epoch 8/20
60000/60000 [=====] - 7s 120us/step - loss: 0.1457
- acc: 0.9576 - val_loss: 0.1517 - val_acc: 0.9530
Epoch 9/20
60000/60000 [=====] - 7s 121us/step - loss: 0.1302
- acc: 0.9618 - val_loss: 0.1327 - val_acc: 0.9594
Epoch 10/20
60000/60000 [=====] - 7s 120us/step - loss: 0.1158
- acc: 0.9662 - val_loss: 0.1224 - val_acc: 0.9617
Epoch 11/20
60000/60000 [=====] - 7s 119us/step - loss: 0.1030
- acc: 0.9694 - val_loss: 0.1135 - val_acc: 0.9646
Epoch 12/20
60000/60000 [=====] - 7s 121us/step - loss: 0.0939
- acc: 0.9722 - val_loss: 0.1117 - val_acc: 0.9674
Epoch 13/20
60000/60000 [=====] - 7s 122us/step - loss: 0.0855
- acc: 0.9742 - val_loss: 0.1008 - val_acc: 0.9694
Epoch 14/20
60000/60000 [=====] - 7s 120us/step - loss: 0.0766
- acc: 0.9769 - val_loss: 0.0925 - val_acc: 0.9710
Epoch 15/20
60000/60000 [=====] - 7s 118us/step - loss: 0.0706
- acc: 0.9788 - val_loss: 0.0936 - val_acc: 0.9711
Epoch 16/20
60000/60000 [=====] - 7s 119us/step - loss: 0.0643
- acc: 0.9801 - val_loss: 0.0921 - val_acc: 0.9723
Epoch 17/20
60000/60000 [=====] - 7s 120us/step - loss: 0.0590
- acc: 0.9818 - val_loss: 0.0844 - val_acc: 0.9746
Epoch 18/20
60000/60000 [=====] - 7s 118us/step - loss: 0.0535
- acc: 0.9841 - val_loss: 0.0808 - val_acc: 0.9746
Epoch 19/20
60000/60000 [=====] - 7s 119us/step - loss: 0.0490
- acc: 0.9853 - val_loss: 0.0783 - val_acc: 0.9761
Epoch 20/20
60000/60000 [=====] - 7s 120us/step - loss: 0.0439
- acc: 0.9872 - val_loss: 0.0764 - val_acc: 0.9759
```

In [9]:

```
# Evaluating the model
score = model_2.evaluate(X_test, y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

# Test and train accuracy of the model
model_2_test = score[1]
model_2_train = history_2.history['acc']


# Plotting Train and Test Loss VS no. of epochs
# List of epoch numbers
x = list(range(1, nb_epoch+1))

# Validation Loss
vy = history_2.history['val_loss']
# Training Loss
ty = history_2.history['loss']

# Calling the function to draw the plot
plt_dynamic(x, vy, ty)
```

Test score: 0.07642286013420671

Test accuracy: 0.9759



(1.b) With dropout and Batch Normalization

In [10]:

```

from keras.layers.normalization import BatchNormalization
from keras.layers import Dropout

# Initialising model
model_2d = Sequential()

# Adding first hidden layer
model_2d.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=he_
# Adding Batch Normalization
model_2d.add(BatchNormalization())
# Adding dropout to first hidden layer
model_2d.add(Dropout(0.4))

# Adding second hidden layer
model_2d.add(Dense(256, activation='relu', kernel_initializer=he_normal(seed=None)))
# Adding Batch Normalization
model_2d.add(BatchNormalization())
# Adding dropout to second hidden layer
model_2d.add(Dropout(0.4))

# Adding output layer
model_2d.add(Dense(output_dim, activation='softmax'))

# Printing model Summary
print("Model Summary :- \n", model_2d.summary())

# Compiling the model
model_2d.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

# Fitting the data to the model
history_2d = model_2d.fit(X_train, y_train, batch_size=batch_size, epochs=nb_epoch, verbose=

```

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:3445: calling dropout (from tensorflow.python.ops.nn_ops) with keep_prob is deprecated and will be removed in a future version. Instructions for updating:
Please use `rate` instead of `keep_prob`. Rate should be set to `rate = 1 - keep_prob`.

Layer (type)	Output Shape	Param #
dense_4 (Dense)	(None, 512)	401920
batch_normalization_1 (Batch Normalization)	(None, 512)	2048
dropout_1 (Dropout)	(None, 512)	0
dense_5 (Dense)	(None, 256)	131328
batch_normalization_2 (Batch Normalization)	(None, 256)	1024
dropout_2 (Dropout)	(None, 256)	0
dense_6 (Dense)	(None, 10)	2570
Total params: 538,890		
Trainable params: 537,354		
Non-trainable params: 1,536		

Model Summary :-

None

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 14s 229us/step - loss: 0.3672
- acc: 0.8881 - val_loss: 0.1736 - val_acc: 0.9480

Epoch 2/20

60000/60000 [=====] - 13s 218us/step - loss: 0.2139
- acc: 0.9354 - val_loss: 0.1216 - val_acc: 0.9627

Epoch 3/20

60000/60000 [=====] - 12s 203us/step - loss: 0.1791
- acc: 0.9451 - val_loss: 0.1077 - val_acc: 0.9673

Epoch 4/20

60000/60000 [=====] - 12s 203us/step - loss: 0.1575
- acc: 0.9523 - val_loss: 0.1018 - val_acc: 0.9676

Epoch 5/20

60000/60000 [=====] - 12s 201us/step - loss: 0.1469
- acc: 0.9544 - val_loss: 0.0884 - val_acc: 0.9717

Epoch 6/20

60000/60000 [=====] - 12s 202us/step - loss: 0.1414
- acc: 0.9567 - val_loss: 0.1041 - val_acc: 0.9669

Epoch 7/20

60000/60000 [=====] - 12s 200us/step - loss: 0.1328
- acc: 0.9585 - val_loss: 0.0933 - val_acc: 0.9722

Epoch 8/20

60000/60000 [=====] - 12s 194us/step - loss: 0.1296
- acc: 0.9600 - val_loss: 0.0903 - val_acc: 0.9726

Epoch 9/20

60000/60000 [=====] - 11s 187us/step - loss: 0.1275
- acc: 0.9601 - val_loss: 0.0869 - val_acc: 0.9738

Epoch 10/20

60000/60000 [=====] - 12s 201us/step - loss: 0.1262
- acc: 0.9614 - val_loss: 0.0755 - val_acc: 0.9762

Epoch 11/20

60000/60000 [=====] - 12s 198us/step - loss: 0.1248
- acc: 0.9611 - val_loss: 0.0790 - val_acc: 0.9753

Epoch 12/20

60000/60000 [=====] - 11s 190us/step - loss: 0.1143
- acc: 0.9640 - val_loss: 0.0789 - val_acc: 0.9771

Epoch 13/20

60000/60000 [=====] - 11s 191us/step - loss: 0.1137
- acc: 0.9650 - val_loss: 0.0746 - val_acc: 0.9760

Epoch 14/20

60000/60000 [=====] - 11s 182us/step - loss: 0.1100
- acc: 0.9658 - val_loss: 0.0697 - val_acc: 0.9797

Epoch 15/20

60000/60000 [=====] - 11s 186us/step - loss: 0.1105
- acc: 0.9648 - val_loss: 0.0701 - val_acc: 0.9784

Epoch 16/20

60000/60000 [=====] - 12s 192us/step - loss: 0.1079
- acc: 0.9658 - val_loss: 0.0747 - val_acc: 0.9775

Epoch 17/20

60000/60000 [=====] - 11s 189us/step - loss: 0.1043
- acc: 0.9676 - val_loss: 0.0722 - val_acc: 0.9788

Epoch 18/20

60000/60000 [=====] - 11s 184us/step - loss: 0.1027
- acc: 0.9674 - val_loss: 0.0730 - val_acc: 0.9788

Epoch 19/20

60000/60000 [=====] - 12s 195us/step - loss: 0.1037
- acc: 0.9678 - val_loss: 0.0625 - val_acc: 0.9804

Epoch 20/20

60000/60000 [=====] - 11s 191us/step - loss: 0.0995
 - acc: 0.9690 - val_loss: 0.0767 - val_acc: 0.9764

In [12]:

```
# Evaluating the model
score = model_2d.evaluate(X_test, y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

# Test and train accuracy of the model
model_2d_test = score[1]
model_2d_train = history_2d.history['acc']

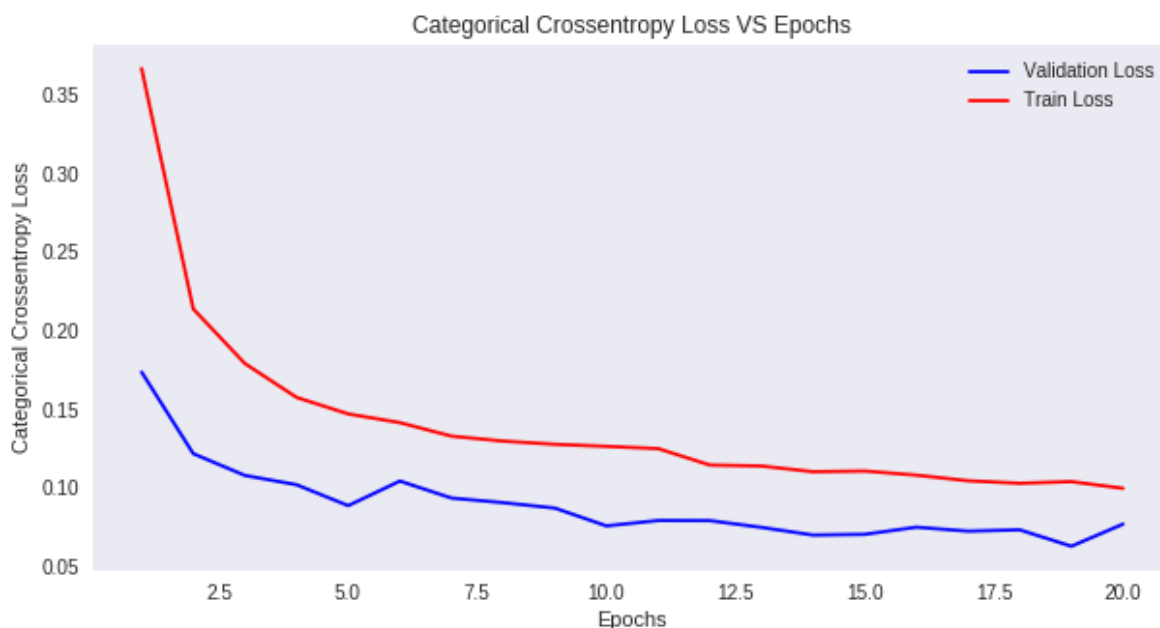
# Plotting Train and Test Loss VS no. of epochs
# List of epoch numbers
x = list(range(1, nb_epoch+1))

# Validation loss
vy = history_2d.history['val_loss']
# Training loss
ty = history_2d.history['loss']

# Calling the function to draw the plot
plt_dynamic(x, vy, ty)
```

Test score: 0.07667584380609915

Test accuracy: 0.9764



(2). Softmax Classifier with 3 hidden layers

(2.a) Without Dropout and Batch Normalization

In [13]:

```

# Initialising model
model_3 = Sequential()

# Adding first hidden layer
model_3.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=he_r

# Adding second hidden layer
model_3.add(Dense(256, activation='relu', kernel_initializer=he_normal(seed=None)))

# Adding third hidden layer
model_3.add(Dense(128, activation='relu', kernel_initializer=he_normal(seed=None)))

# Adding output layer
model_3.add(Dense(output_dim, activation='softmax'))

# Printing model Summary
print(model_3.summary())

# Compiling the model
model_3.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

# Fitting the data to the model
history_3 = model_3.fit(X_train, y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1

```

Layer (type)	Output Shape	Param #
dense_7 (Dense)	(None, 512)	401920
dense_8 (Dense)	(None, 256)	131328
dense_9 (Dense)	(None, 128)	32896
dense_10 (Dense)	(None, 10)	1290
Total params: 567,434		
Trainable params: 567,434		
Non-trainable params: 0		

None

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 11s 187us/step - loss: 0.7319

- acc: 0.7790 - val_loss: 0.3602 - val_acc: 0.8931

Epoch 2/20

60000/60000 [=====] - 11s 179us/step - loss: 0.3282

- acc: 0.9025 - val_loss: 0.2890 - val_acc: 0.9137

Epoch 3/20

60000/60000 [=====] - 11s 181us/step - loss: 0.2564

- acc: 0.9239 - val_loss: 0.2298 - val_acc: 0.9279

Epoch 4/20

60000/60000 [=====] - 10s 172us/step - loss: 0.2048

- acc: 0.9384 - val_loss: 0.1779 - val_acc: 0.9459

Epoch 5/20

60000/60000 [=====] - 11s 182us/step - loss: 0.1654

- acc: 0.9505 - val_loss: 0.1503 - val_acc: 0.9540

Epoch 6/20

60000/60000 [=====] - 10s 172us/step - loss: 0.1374

- acc: 0.9585 - val_loss: 0.1259 - val_acc: 0.9617

```
Epoch 7/20
60000/60000 [=====] - 10s 171us/step - loss: 0.1156
- acc: 0.9656 - val_loss: 0.1206 - val_acc: 0.9635
Epoch 8/20
60000/60000 [=====] - 10s 171us/step - loss: 0.1000
- acc: 0.9700 - val_loss: 0.1032 - val_acc: 0.9676
Epoch 9/20
60000/60000 [=====] - 10s 172us/step - loss: 0.0872
- acc: 0.9736 - val_loss: 0.0944 - val_acc: 0.9712
Epoch 10/20
60000/60000 [=====] - 10s 172us/step - loss: 0.0761
- acc: 0.9765 - val_loss: 0.0901 - val_acc: 0.9730
Epoch 11/20
60000/60000 [=====] - 10s 169us/step - loss: 0.0677
- acc: 0.9793 - val_loss: 0.0842 - val_acc: 0.9731
Epoch 12/20
60000/60000 [=====] - 10s 170us/step - loss: 0.0592
- acc: 0.9817 - val_loss: 0.0806 - val_acc: 0.9756
Epoch 13/20
60000/60000 [=====] - 10s 170us/step - loss: 0.0523
- acc: 0.9838 - val_loss: 0.0779 - val_acc: 0.9756
Epoch 14/20
60000/60000 [=====] - 10s 171us/step - loss: 0.0460
- acc: 0.9857 - val_loss: 0.0816 - val_acc: 0.9754
Epoch 15/20
60000/60000 [=====] - 11s 176us/step - loss: 0.0423
- acc: 0.9863 - val_loss: 0.0786 - val_acc: 0.9768
Epoch 16/20
60000/60000 [=====] - 11s 178us/step - loss: 0.0370
- acc: 0.9881 - val_loss: 0.0816 - val_acc: 0.9750
Epoch 17/20
60000/60000 [=====] - 11s 175us/step - loss: 0.0312
- acc: 0.9906 - val_loss: 0.0796 - val_acc: 0.9764
Epoch 18/20
60000/60000 [=====] - 11s 178us/step - loss: 0.0288
- acc: 0.9907 - val_loss: 0.0746 - val_acc: 0.9774
Epoch 19/20
60000/60000 [=====] - 11s 179us/step - loss: 0.0275
- acc: 0.9913 - val_loss: 0.0851 - val_acc: 0.9747
Epoch 20/20
60000/60000 [=====] - 11s 180us/step - loss: 0.0237
- acc: 0.9925 - val_loss: 0.0804 - val_acc: 0.9771
```

In [14]:

```
# Evaluating the model
score = model_3.evaluate(X_test, y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

# Test and train accuracy of the model
model_3_test = score[1]
model_3_train = history_3.history['acc']


# Plotting Train and Test Loss VS no. of epochs
# List of epoch numbers
x = list(range(1,nb_epoch+1))

# Validation loss
vy = history_3.history['val_loss']
# Training loss
ty = history_3.history['loss']

# Calling the function to draw the plot
plt_dynamic(x, vy, ty)
```

Test score: 0.08041134356782713

Test accuracy: 0.9771



(2.b) With Dropout and Batch Normalization

In [15]:

```

model_3d = Sequential()

# Adding first hidden layer
model_3d.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=he_
# Adding Batch Normalization
model_3d.add(BatchNormalization())
# Adding dropout
model_3d.add(Dropout(0.4))

# Adding second hidden layer
model_3d.add(Dense(256, activation='relu', kernel_initializer=he_normal(seed=None)))
# Adding Batch Normalization
model_3d.add(BatchNormalization())
# Adding dropout
model_3d.add(Dropout(0.4))

# Adding third hidden layer
model_3d.add(Dense(128, activation='relu', kernel_initializer=he_normal(seed=None)))
# Adding Batch Normalization
model_3d.add(BatchNormalization())
# Adding dropout
model_3d.add(Dropout(0.4))

# Adding output layer
model_3d.add(Dense(output_dim, activation='softmax'))

# Printing model Summary
print(model_3d.summary())

# Compiling the model
model_3d.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

# Fitting the data to the model
history_3d = model_3d.fit(X_train, y_train, batch_size=batch_size, epochs=nb_epoch, verbose=

```

Layer (type)	Output Shape	Param #
=====		
dense_11 (Dense)	(None, 512)	401920
batch_normalization_3 (Batch Normalization)	(None, 512)	2048
dropout_3 (Dropout)	(None, 512)	0
dense_12 (Dense)	(None, 256)	131328
batch_normalization_4 (Batch Normalization)	(None, 256)	1024
dropout_4 (Dropout)	(None, 256)	0
dense_13 (Dense)	(None, 128)	32896
batch_normalization_5 (Batch Normalization)	(None, 128)	512
dropout_5 (Dropout)	(None, 128)	0
dense_14 (Dense)	(None, 10)	1290
=====		

Total params: 571,018
Trainable params: 569,226
Non-trainable params: 1,792

None

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 13s 225us/step - loss: 0.4926
- acc: 0.8490 - val_loss: 0.1797 - val_acc: 0.9451

Epoch 2/20

60000/60000 [=====] - 12s 198us/step - loss: 0.2733
- acc: 0.9176 - val_loss: 0.1389 - val_acc: 0.9561

Epoch 3/20

60000/60000 [=====] - 12s 200us/step - loss: 0.2253
- acc: 0.9319 - val_loss: 0.1387 - val_acc: 0.9557

Epoch 4/20

60000/60000 [=====] - 12s 207us/step - loss: 0.1986
- acc: 0.9391 - val_loss: 0.1084 - val_acc: 0.9660

Epoch 5/20

60000/60000 [=====] - 12s 200us/step - loss: 0.1800
- acc: 0.9454 - val_loss: 0.0997 - val_acc: 0.9699

Epoch 6/20

60000/60000 [=====] - 13s 216us/step - loss: 0.1655
- acc: 0.9482 - val_loss: 0.0973 - val_acc: 0.9691

Epoch 7/20

60000/60000 [=====] - 13s 210us/step - loss: 0.1628
- acc: 0.9503 - val_loss: 0.0925 - val_acc: 0.9698

Epoch 8/20

60000/60000 [=====] - 12s 205us/step - loss: 0.1466
- acc: 0.9549 - val_loss: 0.0996 - val_acc: 0.9666

Epoch 9/20

60000/60000 [=====] - 12s 203us/step - loss: 0.1444
- acc: 0.9556 - val_loss: 0.0894 - val_acc: 0.9720

Epoch 10/20

60000/60000 [=====] - 12s 200us/step - loss: 0.1391
- acc: 0.9566 - val_loss: 0.0793 - val_acc: 0.9767

Epoch 11/20

60000/60000 [=====] - 12s 201us/step - loss: 0.1343
- acc: 0.9587 - val_loss: 0.0773 - val_acc: 0.9756

Epoch 12/20

60000/60000 [=====] - 12s 195us/step - loss: 0.1340
- acc: 0.9587 - val_loss: 0.0749 - val_acc: 0.9761

Epoch 13/20

60000/60000 [=====] - 12s 201us/step - loss: 0.1280
- acc: 0.9600 - val_loss: 0.0821 - val_acc: 0.9749

Epoch 14/20

60000/60000 [=====] - 12s 193us/step - loss: 0.1261
- acc: 0.9607 - val_loss: 0.0772 - val_acc: 0.9744

Epoch 15/20

60000/60000 [=====] - 12s 204us/step - loss: 0.1250
- acc: 0.9608 - val_loss: 0.0709 - val_acc: 0.9768

Epoch 16/20

60000/60000 [=====] - 12s 208us/step - loss: 0.1201
- acc: 0.9619 - val_loss: 0.0715 - val_acc: 0.9776

Epoch 17/20

60000/60000 [=====] - 12s 207us/step - loss: 0.1176
- acc: 0.9638 - val_loss: 0.0671 - val_acc: 0.9794

Epoch 18/20

60000/60000 [=====] - 12s 208us/step - loss: 0.1152
- acc: 0.9642 - val_loss: 0.0665 - val_acc: 0.9785

Epoch 19/20

```
60000/60000 [=====] - 12s 205us/step - loss: 0.1115
- acc: 0.9649 - val_loss: 0.0692 - val_acc: 0.9774
Epoch 20/20
60000/60000 [=====] - 13s 209us/step - loss: 0.1134
- acc: 0.9644 - val_loss: 0.0695 - val_acc: 0.9780
```

In [16]:

```
# Evaluating the model
score = model_3d.evaluate(X_test, y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

# Test and train accuracy of the model
model_3d_test = score[1]
model_3d_train = history_3d.history['acc']


# Plotting Train and Test Loss VS no. of epochs
# List of epoch numbers
x = list(range(1, nb_epoch+1))

# Validation Loss
vy = history_3d.history['val_loss']
# Training Loss
ty = history_3d.history['loss']

# Calling the function to draw the plot
plt_dynamic(x, vy, ty)
```

Test score: 0.06945714832819067

Test accuracy: 0.978



(3). Softmax Classifier with 5 hidden layers

(3.a) Without Dropout and Batch Normalization

In [17]:

```

# Initialising model
model_5 = Sequential()

# Adding first hidden layer
model_5.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=he_r

# Adding second hidden layer
model_5.add(Dense(256, activation='relu', kernel_initializer=he_normal(seed=None)))

# Adding third hidden layer
model_5.add(Dense(128, activation='relu', kernel_initializer=he_normal(seed=None)))

# Adding fourth hidden layer
model_5.add(Dense(64, activation='relu', kernel_initializer=he_normal(seed=None)))

# Adding fifth hidden layer
model_5.add(Dense(32, activation='relu', kernel_initializer=he_normal(seed=None)))

# Adding output layer
model_5.add(Dense(output_dim, activation='softmax'))

# Printing model Summary
print(model_5.summary())

# Compiling the model
model_5.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

# Fitting the data to the model
history_5 = model_5.fit(X_train, y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1

```

Layer (type)	Output Shape	Param #
dense_15 (Dense)	(None, 512)	401920
dense_16 (Dense)	(None, 256)	131328
dense_17 (Dense)	(None, 128)	32896
dense_18 (Dense)	(None, 64)	8256
dense_19 (Dense)	(None, 32)	2080
dense_20 (Dense)	(None, 10)	330
Total params: 576,810		
Trainable params: 576,810		
Non-trainable params: 0		

None

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 12s 192us/step - loss: 0.8731
 - acc: 0.7113 - val_loss: 0.4438 - val_acc: 0.8727

Epoch 2/20

60000/60000 [=====] - 11s 180us/step - loss: 0.3234
 - acc: 0.9059 - val_loss: 0.2611 - val_acc: 0.9229

Epoch 3/20

60000/60000 [=====] - 11s 182us/step - loss: 0.2123

```
- acc: 0.9378 - val_loss: 0.1829 - val_acc: 0.9456
Epoch 4/20
60000/60000 [=====] - 11s 180us/step - loss: 0.1610
- acc: 0.9521 - val_loss: 0.1379 - val_acc: 0.9589
Epoch 5/20
60000/60000 [=====] - 11s 176us/step - loss: 0.1254
- acc: 0.9620 - val_loss: 0.1246 - val_acc: 0.9621
Epoch 6/20
60000/60000 [=====] - 11s 178us/step - loss: 0.1055
- acc: 0.9674 - val_loss: 0.1105 - val_acc: 0.9659
Epoch 7/20
60000/60000 [=====] - 11s 181us/step - loss: 0.0903
- acc: 0.9731 - val_loss: 0.1066 - val_acc: 0.9695
Epoch 8/20
60000/60000 [=====] - 10s 171us/step - loss: 0.0793
- acc: 0.9753 - val_loss: 0.1072 - val_acc: 0.9662
Epoch 9/20
60000/60000 [=====] - 11s 184us/step - loss: 0.0716
- acc: 0.9778 - val_loss: 0.1158 - val_acc: 0.9657
Epoch 10/20
60000/60000 [=====] - 11s 186us/step - loss: 0.0606
- acc: 0.9812 - val_loss: 0.1062 - val_acc: 0.9683
Epoch 11/20
60000/60000 [=====] - 11s 184us/step - loss: 0.0536
- acc: 0.9832 - val_loss: 0.0847 - val_acc: 0.9732
Epoch 12/20
60000/60000 [=====] - 11s 185us/step - loss: 0.0490
- acc: 0.9845 - val_loss: 0.0822 - val_acc: 0.9765
Epoch 13/20
60000/60000 [=====] - 11s 185us/step - loss: 0.0429
- acc: 0.9856 - val_loss: 0.0879 - val_acc: 0.9756
Epoch 14/20
60000/60000 [=====] - 11s 175us/step - loss: 0.0411
- acc: 0.9865 - val_loss: 0.0967 - val_acc: 0.9753
Epoch 15/20
60000/60000 [=====] - 11s 176us/step - loss: 0.0353
- acc: 0.9888 - val_loss: 0.0875 - val_acc: 0.9766
Epoch 16/20
60000/60000 [=====] - 11s 177us/step - loss: 0.0334
- acc: 0.9892 - val_loss: 0.0915 - val_acc: 0.9761
Epoch 17/20
60000/60000 [=====] - 11s 183us/step - loss: 0.0295
- acc: 0.9900 - val_loss: 0.0871 - val_acc: 0.9770
Epoch 18/20
60000/60000 [=====] - 10s 169us/step - loss: 0.0280
- acc: 0.9908 - val_loss: 0.0865 - val_acc: 0.9769
Epoch 19/20
60000/60000 [=====] - 10s 167us/step - loss: 0.0288
- acc: 0.9907 - val_loss: 0.1064 - val_acc: 0.9729
Epoch 20/20
60000/60000 [=====] - 10s 161us/step - loss: 0.0233
- acc: 0.9921 - val_loss: 0.0884 - val_acc: 0.9783
```

In [19]:

```
# Evaluating the model
score = model_5.evaluate(X_test, y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

# Test and train accuracy of the model
model_5_test = score[1]
model_5_train = history_5.history['acc']


# Plotting Train and Test Loss VS no. of epochs
# List of epoch numbers
x = list(range(1, nb_epoch+1))

# Validation Loss
vy = history_5.history['val_loss']
# Training Loss
ty = history_5.history['loss']

# Calling the function to draw the plot
plt_dynamic(x, vy, ty)
```

Test score: 0.08839814674687223

Test accuracy: 0.9783



(3.b) With Dropout and Batch Normalisation

In [20]:

```

# Initialising model
model_5d = Sequential()

# Adding first hidden layer
model_5d.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=he_
# Adding Batch Normalization
model_5d.add(BatchNormalization())
# Adding dropout
model_5d.add(Dropout(0.4))

# Adding second hidden layer
model_5d.add(Dense(256, activation='relu', kernel_initializer=he_normal(seed=None)))
# Adding Batch Normalization
model_5d.add(BatchNormalization())
# Adding dropout
model_5d.add(Dropout(0.4))

# Adding third hidden layer
model_5d.add(Dense(128, activation='relu', kernel_initializer=he_normal(seed=None)))
# Adding Batch Normalization
model_5d.add(BatchNormalization())
# Adding dropout
model_5d.add(Dropout(0.4))

# Adding fourth hidden layer
model_5d.add(Dense(64, activation='relu', kernel_initializer=he_normal(seed=None)))
# Adding Batch Normalization
model_5d.add(BatchNormalization())
# Adding dropout
model_5d.add(Dropout(0.4))

# Adding fifth hidden layer
model_5d.add(Dense(32, activation='relu', kernel_initializer=he_normal(seed=None)))
# Adding Batch Normalization
model_5d.add(BatchNormalization())
# Adding dropout
model_5d.add(Dropout(0.4))

# Adding output layer
model_5d.add(Dense(output_dim, activation='softmax'))

# Printing model Summary
print(model_5d.summary())

# Compiling the model
model_5d.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

# Fitting the data to the model
history_5d = model_5d.fit(X_train, y_train, batch_size=batch_size, epochs=nb_epoch, verbose=

```

Layer (type)	Output Shape	Param #
=====		
dense_21 (Dense)	(None, 512)	401920
batch_normalization_6 (Batch Normalization)	(None, 512)	2048
dropout_6 (Dropout)	(None, 512)	0

dense_22 (Dense)	(None, 256)	131328
batch_normalization_7 (Batch Normalization)	(None, 256)	1024
dropout_7 (Dropout)	(None, 256)	0
dense_23 (Dense)	(None, 128)	32896
batch_normalization_8 (Batch Normalization)	(None, 128)	512
dropout_8 (Dropout)	(None, 128)	0
dense_24 (Dense)	(None, 64)	8256
batch_normalization_9 (Batch Normalization)	(None, 64)	256
dropout_9 (Dropout)	(None, 64)	0
dense_25 (Dense)	(None, 32)	2080
batch_normalization_10 (Batch Normalization)	(None, 32)	128
dropout_10 (Dropout)	(None, 32)	0
dense_26 (Dense)	(None, 10)	330
=====		
Total params: 580,778		
Trainable params: 578,794		
Non-trainable params: 1,984		

None

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 15s 243us/step - loss: 1.0536
- acc: 0.6618 - val_loss: 0.3050 - val_acc: 0.9124

Epoch 2/20

60000/60000 [=====] - 13s 212us/step - loss: 0.4741
- acc: 0.8712 - val_loss: 0.2487 - val_acc: 0.9302

Epoch 3/20

60000/60000 [=====] - 13s 223us/step - loss: 0.3790
- acc: 0.9002 - val_loss: 0.2027 - val_acc: 0.9430

Epoch 4/20

60000/60000 [=====] - 13s 222us/step - loss: 0.3324
- acc: 0.9130 - val_loss: 0.1723 - val_acc: 0.9490

Epoch 5/20

60000/60000 [=====] - 13s 221us/step - loss: 0.2972
- acc: 0.9239 - val_loss: 0.1530 - val_acc: 0.9568

Epoch 6/20

60000/60000 [=====] - 13s 221us/step - loss: 0.2780
- acc: 0.9283 - val_loss: 0.1594 - val_acc: 0.9556

Epoch 7/20

60000/60000 [=====] - 13s 213us/step - loss: 0.2545
- acc: 0.9347 - val_loss: 0.1426 - val_acc: 0.9602

Epoch 8/20

60000/60000 [=====] - 13s 221us/step - loss: 0.2501
- acc: 0.9358 - val_loss: 0.1360 - val_acc: 0.9595

Epoch 9/20

60000/60000 [=====] - 13s 218us/step - loss: 0.2390
- acc: 0.9388 - val_loss: 0.1300 - val_acc: 0.9611

Epoch 10/20

60000/60000 [=====] - 13s 218us/step - loss: 0.2301

```
- acc: 0.9408 - val_loss: 0.1213 - val_acc: 0.9656
Epoch 11/20
60000/60000 [=====] - 13s 217us/step - loss: 0.2175
- acc: 0.9436 - val_loss: 0.1177 - val_acc: 0.9677
Epoch 12/20
60000/60000 [=====] - 13s 219us/step - loss: 0.2136
- acc: 0.9445 - val_loss: 0.1163 - val_acc: 0.9678
Epoch 13/20
60000/60000 [=====] - 13s 209us/step - loss: 0.2083
- acc: 0.9459 - val_loss: 0.1036 - val_acc: 0.9727
Epoch 14/20
60000/60000 [=====] - 13s 214us/step - loss: 0.1997
- acc: 0.9482 - val_loss: 0.1060 - val_acc: 0.9703
Epoch 15/20
60000/60000 [=====] - 13s 219us/step - loss: 0.1936
- acc: 0.9494 - val_loss: 0.0975 - val_acc: 0.9743
Epoch 16/20
60000/60000 [=====] - 13s 218us/step - loss: 0.1910
- acc: 0.9509 - val_loss: 0.0997 - val_acc: 0.9729
Epoch 17/20
60000/60000 [=====] - 13s 218us/step - loss: 0.1883
- acc: 0.9506 - val_loss: 0.0982 - val_acc: 0.9730
Epoch 18/20
60000/60000 [=====] - 13s 221us/step - loss: 0.1842
- acc: 0.9521 - val_loss: 0.1073 - val_acc: 0.9694
Epoch 19/20
60000/60000 [=====] - 13s 217us/step - loss: 0.1835
- acc: 0.9526 - val_loss: 0.0940 - val_acc: 0.9739
Epoch 20/20
60000/60000 [=====] - 13s 218us/step - loss: 0.1801
- acc: 0.9535 - val_loss: 0.0928 - val_acc: 0.9751
```

In [21]:

```
# Evaluating the model
score = model_5d.evaluate(X_test, y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

# Test and train accuracy of the model
model_5d_test = score[1]
model_5d_train = history_5d.history['acc']


# Plotting Train and Test Loss VS no. of epochs
# List of epoch numbers
x = list(range(1,nb_epoch+1))

# Validation Loss
vy = history_5d.history['val_loss']
# Training Loss
ty = history_5d.history['loss']

# Calling the function to draw the plot
plt_dynamic(x, vy, ty)
```

Test score: 0.09278820200497284

Test accuracy: 0.9751



CONCLUSION

(a). Procedure Followed :

1. Load MNIST dataset
2. Split the dataset into train and test
3. Normalize the train and test data
4. Convert class variable into categorical data vector
5. Implement Softmax classifier with 2 , 3 and 5 hidden layers(512,256,128,64,32 respectively) .
6. Add Dropout(rate 40%) and Batch Normalization to the hidden layers .

7. Draw Categorical Crossentropy Loss VS No.of Epochs plot .

(b) Table (Different models with their train and test accuracies):

In [22]:

```
from prettytable import PrettyTable

x = PrettyTable()

x.field_names = ["MLP_Hidden_Layers", "MODEL", "Training Accuracy", "Test Accuracy"]

x.add_row([2, "Without Dropout and Batch Normalization", 0.95, 0.95])
x.add_row([2, "With Dropout and Batch Normalization", 0.96, 0.97])
x.add_row([3, "Without Dropout and Batch Normalization", 0.95, 0.95])
x.add_row([3, "With Dropout and Batch Normalization", 0.95, 0.95])
x.add_row([5, "Without Dropout and Batch Normalization", 0.96, 0.96])
x.add_row([5, "With Dropout and Batch Normalization", 0.95, 0.96])

print(x)
```

MLP_Hidden_Layers	MODEL	Training Accuracy	Test Accuracy
2	Without Dropout and Batch Normalization	0.95	0.95
2	With Dropout and Batch Normalization	0.96	0.97
3	Without Dropout and Batch Normalization	0.95	0.95
3	With Dropout and Batch Normalization	0.95	0.95
5	Without Dropout and Batch Normalization	0.96	0.96
5	With Dropout and Batch Normalization	0.95	0.96

In [0]: